

# Homework Assignment #1: I/O Performance under Sequential and Random R/W Workloads

# Outline

- **File Input / Output**
- Application Programming Interface
  - POSIX File I/O
  - Other APIs
- Homework Assignment #1
- Reference

# File Input / Output

- Read Operation

Process A

```
int main()
{
  ...
  read(fd, buffer, 4096);
  ...
  return 0;
}
```

User  
mode

Page Cache

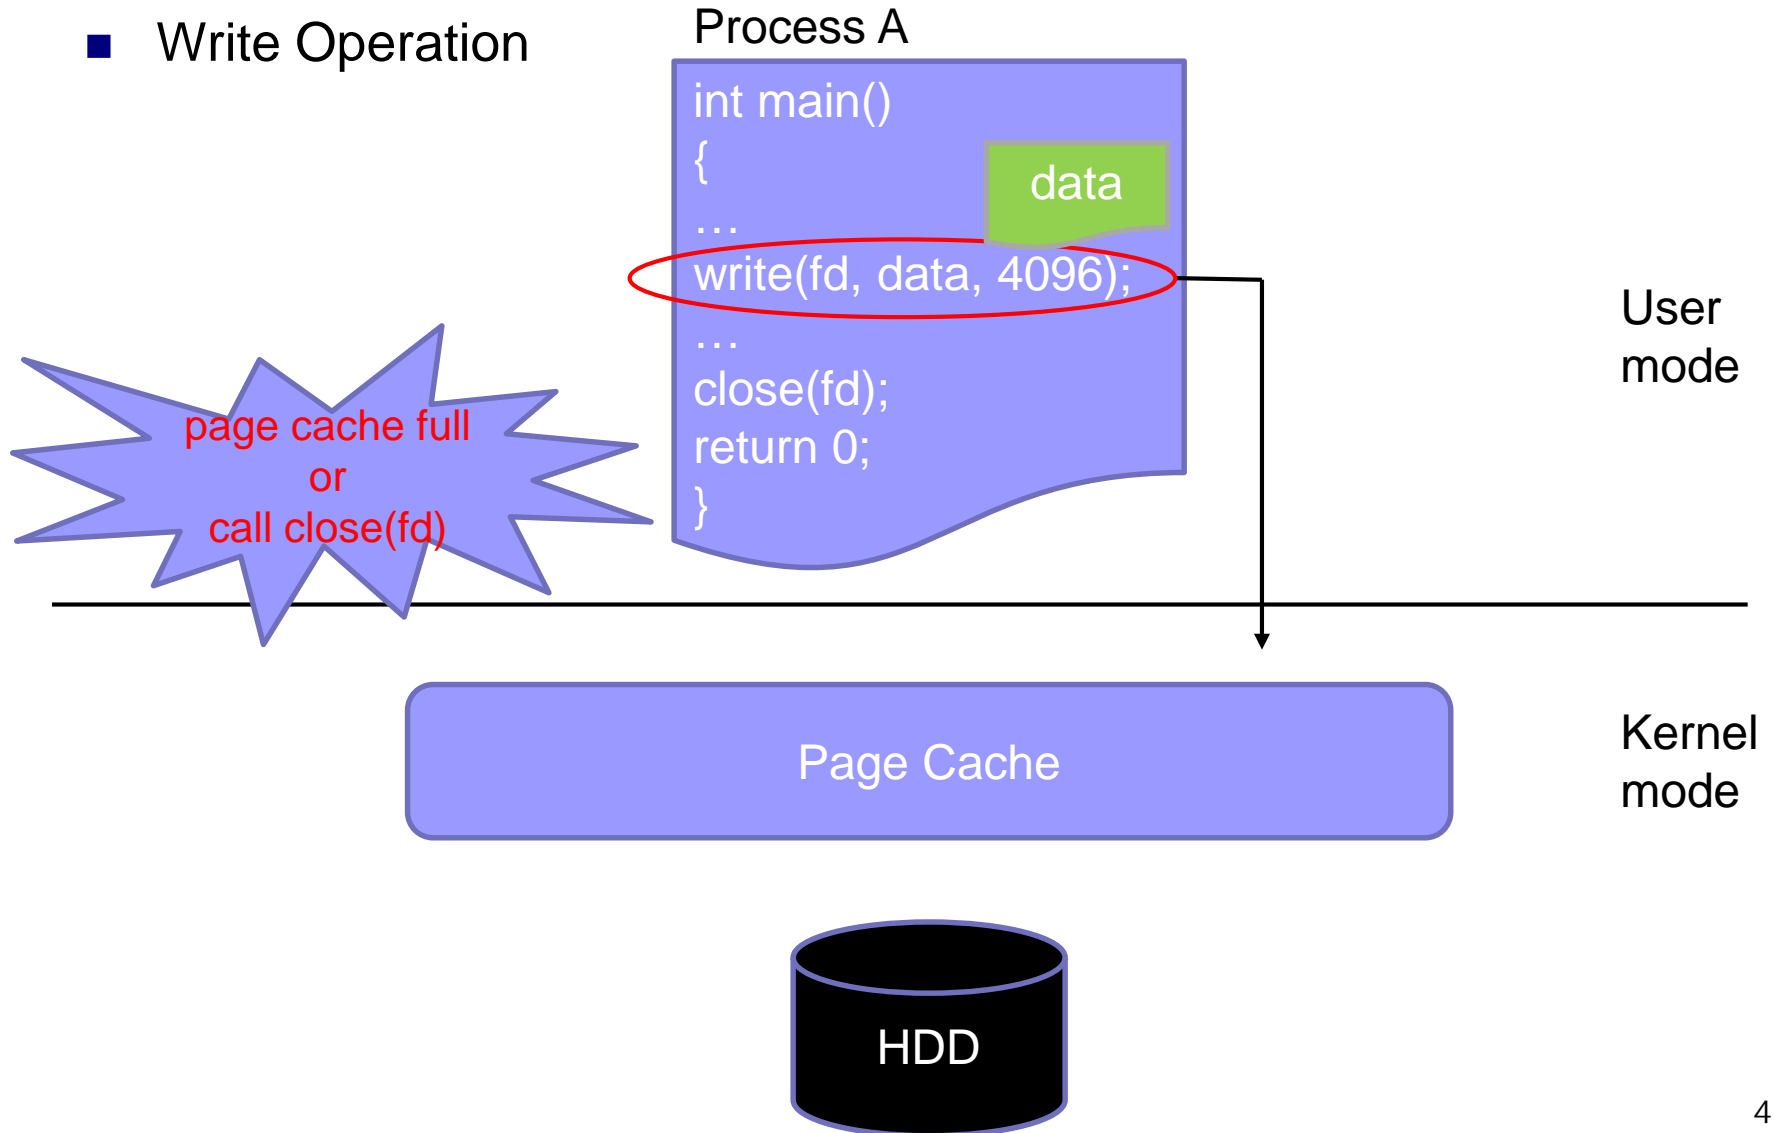
data

Kernel  
mode

block

# File Input / Output

## ■ Write Operation



# Outline

- File Input / Output
- **Application Programming Interface**
  - POSIX File I/O
  - Other APIs
- Homework Assignment #1
- Reference

# POSIX File I/O

- We will use the following POSIX File I/O system calls
  - `int open(const char *pathname, int flags, mode_t mode)`
  - `ssize_t read(int fd, void *buf, size_t count)`
  - `ssize_t write(int fd, const void *buf, size_t count)`
  - `off_t lseek(int fd, off_t offset, int whence)`
  - `int close(int fd);`

# POSIX File I/O (cont.)

## ■ Open File

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
```

```
int open(const char *pathname, int flags, ... /* mode_t mode */);
```

Returns file descriptor on success, or -1 on error

```
ex. : int fd = open( "~/hw1/filename", O_RDONLY, S_IRWXU );
```

**pathname** : file path.

**flags** : The flags argument is a bit mask that specifies the access mode for the file

**mode** : the mode bit-mask argument specifies the permissions to be placed on the file

# POSIX File I/O (cont.)

## ■ Open File (cont.)

```
int open(const char *pathname, int flags, ... /* mode_t mode */);
```

Returns file descriptor on success, or -1 on error

Flag	Purpose	SUS?
O_RDONLY	Open for reading only	v3
O_WRONLY	Open for writing only	v3
O_RDWR	Open for reading and writing	v3
O_CLOEXEC	Set the close-on-exec flag (since Linux 2.6.23)	v4
O_CREAT	Create file if it doesn't already exist	v3
O_DIRECT	File I/O bypasses buffer cache	
O_DIRECTORY	Fail if <i>pathname</i> is not a directory	v4
O_EXCL	With O_CREAT: create file exclusively	v3
O_LARGEFILE	Used on 32-bit systems to open large files	
O_NOATIME	Don't update file last access time on <i>read()</i> (since Linux 2.6.8)	
O_NOCTTY	Don't let <i>pathname</i> become the controlling terminal	v3
O_NOFOLLOW	Don't dereference symbolic links	v4
O_TRUNC	Truncate existing file to zero length	v3
O_APPEND	Writes are always appended to end of file	v3
O_ASYNC	Generate a signal when I/O is possible	
O_DSYNC	Provide synchronized I/O data integrity (since Linux 2.6.33)	v3
O_NONBLOCK	Open in nonblocking mode	v3
O_SYNC	Make file writes synchronous	v3



# POSIX File I/O (cont.)

## ■ Read File

```
#include <unistd.h>
```

```
ssize_t read(int fd, void *buffer, size_t count);
```

Returns number of bytes read, 0 on EOF, or -1 on error

```
ex. : int ret = read( fd, buffer, 4096);
```

***fd*** : file descriptor.

***buffer*** : The buffer argument supplies the address of the memory buffer into which the input data is to be placed.

***count*** : The count argument specifies the maximum number of bytes to read.

# POSIX File I/O (cont.)

## ■ Write File

```
#include <unistd.h>
```

```
ssize_t write(int fd, void *buffer, size_t count);
```

Returns number of bytes written, or -1 on error

```
ex. : int ret = write( fd, newData, 4096);
```

***fd*** : file descriptor.

***buffer*** : buffer is the address of the data to be written.

***count*** : count is the number of bytes to write from buffer

# POSIX File I/O (cont.)

## ■ Changing the File Offset

```
#include <unistd.h>
```

```
off_t lseek(int fd, off_t offset, int whence);
```

Returns new file offset if successful, or -1 on error

```
ex. : int curr = lseek( fd, 0, SEEK_CUR);
```

***fd*** : file descriptor.

***offset*** : The offset argument specifies a value in bytes.

***whence*** : The whence argument indicates the base point from which offset is to be interpreted

# POSIX File I/O (cont.)

## ■ Changing the File Offset (cont.)

```
off_t lseek(int fd, off_t offset, int whence);
```

Returns new file offset if successful, or -1 on error

SEEK\_SET

The file offset is set *offset* bytes from the beginning of the file.

SEEK\_CUR

The file offset is adjusted by *offset* bytes relative to the current file offset.

SEEK\_END

The file offset is set to the size of the file plus *offset*. In other words, *offset* is interpreted with respect to the next byte after the last byte of the file.

# POSIX File I/O (cont.)

## ■ Close File

```
#include <unistd.h>
```

```
int close(int fd);
```

Returns 0 on success, or -1 on error

```
ex. : int ret = close( fd );
```

***fd*** : file descriptor.

# Other APIs

## ■ Allocate Aligned Memory

```
#include <stdlib.h>
```

```
void *valloc(size_t size)
```

return a pointer to the allocated memory, or NULL if the request fails

```
ex. : char *buffer = valloc( sizeof(char) * 4096 );
```

**size** : allocates size bytes.

The memory address will be a multiple of the page size.

# Other APIs (cont.)

## ■ Messure Time

```
#include <sys/time.h>
```

```
int gettimeofday( struct timeval *tv, struct timezone *tz );
```

```
struct timeval {  
    time_t tv_sec;           /* seconds */  
    suseconds_t tv_usec;    /* microseconds */  
};
```

```
struct timezone {  
    int tz_minuteswest;     /* minutes west of Greenwich */  
    int tz_dsttime;         /* type of DST correction */  
};
```

# Other APIs (cont.)

- Measure Time (cont.)

```
int main()
{
    struct timeval start;
    struct timeval end;

    unsigned long diff;
    gettimeofday(&start, NULL);
    delay(10);
    gettimeofday(&end, NULL);
    diff = 1000000 * (end.tv_sec - start.tv_sec) +
           end.tv_usec - start.tv_usec;
    printf("the difference is %ld (us)\n", diff);
    return 0;
}
```



# Other APIs (cont.)

## ■ rand()

- rand() function is provided by stdlib.h library.
- Returns a pseudo-random number in the range of 0 to RAND\_MAX.

```
#include<stdlib.h>
```

```
int rand(void);
```

```
ex. : rand();
```

# Other APIs (cont.)

## ■ **srand()**

- `srand()` function is provided by `stdlib.h` library.
- This function seeds the random number generator used by the function `rand()`.

```
#include<stdlib.h>
```

```
void srand(unsigned int seed);
```

```
ex. : srand(1000);
```

# Other APIs (cont.)

## ■ **time()**

- time() function is provided by time.h library.
- Calculates the current calendar time and encodes it into time\_t format.

```
#include<time.h>
```

```
time_t time(time_t *t);
```

```
ex. : time(NULL);
```

# Example

```
5 #include <stdio.h>
6 #include <stdlib.h>
7 #include <sys/types.h> //open()
8 #include <sys/stat.h> //open()
9 #include <fcntl.h> //open()
10 #include <unistd.h> //read()
11 #include <sys/time.h> //gettimeofday()
12 #include <stdint.h>
13
14 int main()
15 {
16     int32_t fd;
17     int8_t *buffer = valloc(sizeof(char) * 4096);
18     int32_t ret;
19     int64_t exTime;
20     struct timeval timeStart;
21     struct timeval timeEnd;
22
23     fd = open("data", O_RDONLY);
24     if (fd == -1){
25         printf("open error\n");
26         return 0;
27     }
28     gettimeofday(&timeStart, NULL);
29     ret = read(fd, buffer, 4096);
30     gettimeofday(&timeEnd, NULL);
31     if (ret == -1){
32         printf("file read error\n");
33         return 0;
34     }
35     free(buffer);
36     close(fd);
37
38     exTime = 1000000 * (timeEnd.tv_sec - timeStart.tv_sec) + (timeEnd.tv_usec - timeStart.tv_usec);
39
40     printf("exTime = %ld us\n", exTime);
41     return 0;
42 }
```

# Random number example

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
void main()
{
    int a;
    srand(time(NULL));
    for(int i=0; i<=5; i++){
        a=(rand()%100)+1;
        printf("The Random Number is %d \n", a);
    }
}
```

```
arthur@arthur-VirtualBox:~/hw2$ ./rand
The Random Number is 14
The Random Number is 87
The Random Number is 7
The Random Number is 27
The Random Number is 98
The Random Number is 37
```

# Outline

- File Input / Output
- Application Programming Interface
  - POSIX File I/O
  - Other APIs
- **Homework Assignment #1**
- Reference

# Homework Assignments #1: Sequential

- **1st:** Write a program that creates a 100MB file on your local disk
- **2nd:** Measure the time of following three I/O access scenarios
- **Sequential Read / Write (30%)**
  - Sequential Read:
    - Read the file sequentially by reading the file from beginning to end.
  - Sequential Write:
    - Overwrite the file with 100MB of new data by writing the file from beginning to end.
  - Sequential Write (using O\_DIRECT):
    - Overwrite the file with 100MB of new data by writing the file using O\_DIRECT flag from beginning to end.

# Homework Assignments #1: Random

- **1st:** Write a program that creates a 100MB file on your local disk
- **2nd:** Measure the time of following three I/O access scenarios
- Random Read / Write(Do 50,000 times) (30%)
  - Random Read:
    - Choose 4KB-aligned offset in the file uniformly at random, seek to that location in the file, and read 4KB of data at that position.
  - Random Write:
    - Choose 4KB-aligned offset in the file uniformly at random, seek to that location in the file, and overwrite the file with 4KB of new data at that position.
  - Random Write (using O\_DIRECT):
    - Choose 4KB-aligned offset in the file uniformly at random, seek to that location in the file, and overwrite the file using O\_DIRECT flag with 4KB of new data at that position.



# Homework Assignments #1

- Sequential Read / Write (30%)
  - each program 10%
- Random Read / Write (30%)
  - each program 10%
- Word (40%)
  - Explain your execution times.
    1. Compare and explain your results of Sequential Read, Sequential Write and (O\_DIRECT) Sequential Write. (8%)
    2. Compare and explain your results of Random Read, Random Write and (O\_DIRECT) Random Write. (8%)
    3. Compare and explain your results of Sequential Read and Random Read. (8%)
    4. Compare and explain your results of Sequential Write and Random Write. (8%)
    5. Compare and explain your results of (O\_DIRECT) Sequential Write and (O\_DIRECT) Random Write. (8%)

# Appendix

- Get readahead size

```
arthur@arthur-VirtualBox:~$ sudo /sbin/blockdev --getra /dev/sda  
256
```

- Set readahead to 0

```
arthur@arthur-VirtualBox:~$ sudo /sbin/blockdev --setra 0 /dev/sda
```

- Clear page cache, dentries, inode

- ☐ 1st: Sync dirty data to disk

```
arthur@arthur-VirtualBox:~$ sudo sync
```

- ☐ 2nd: clear page cache

```
arthur@arthur-VirtualBox:~$ sudo sh -c 'echo 3 > /proc/sys/vm/drop_caches'
```

# Appendix

- How to use *O\_DIRECT* flag

```
fd = open("data", O_WRONLY | O_DIRECT);
```

- Include header

```
#include <sys/types.h> //open()
#include <sys/stat.h>   //open()
#define __USE_GNU 1
#include <fcntl.h>      //open()
```

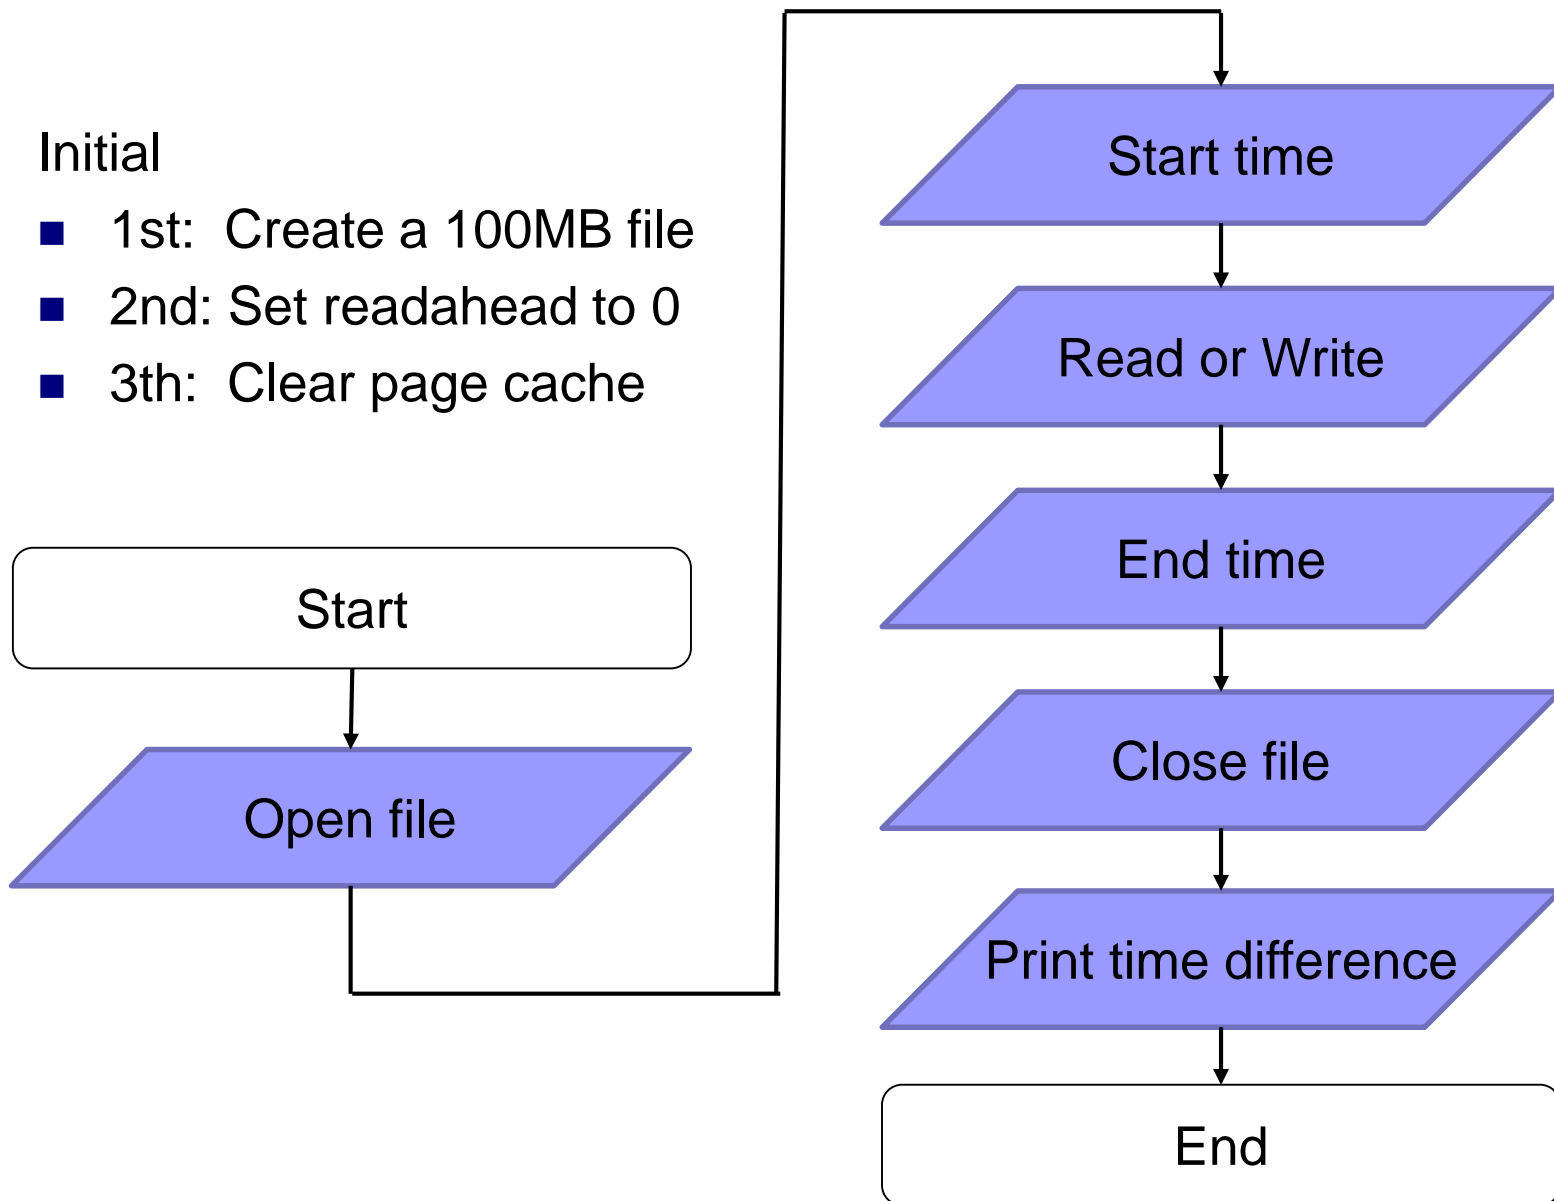
- Add *#define \_\_USE\_GNU 1* over *#include <fcntl.h>*

```
#define __USE_GNU 1
#include <fcntl.h> //open()
```

# Flowchart

## Initial

- 1st: Create a 100MB file
- 2nd: Set readahead to 0
- 3th: Clear page cache



# Result

- Sequential Read

```
oslab@oslab-ASUSPRO-D840MB-M840MB:~/filesystem_hw1/hw1/seqRead$ ./seqRead.out  
seqRead exTime = 1395255 us
```

- Sequential Write

```
oslab@oslab-ASUSPRO-D840MB-M840MB:~/filesystem_hw1/hw1/seqWrite$ ./seqWrite.out  
seqWrite exTime = 29613 us
```

- Sequential Write (O\_DIRECT)

```
oslab@oslab-ASUSPRO-D840MB-M840MB:~/filesystem_hw1/hw1/seqWrite_direct$ ./seqWrite_direct.out  
seqWrite_direct exTime = 505052 us
```

# Result

- Random Read

```
oslab@oslab-ASUSPRO-D840MB-M840MB:~/filesystem_hw1/hw1/ranRead$ ./ranRead.out  
ranRead exTime = 87206953 us
```

- Random Write

```
oslab@oslab-ASUSPRO-D840MB-M840MB:~/filesystem_hw1/hw1/ranWrite$ ./ranWrite.out  
ranWrite exTime = 60447 us
```

- Random Write (O\_DIRECT)

```
oslab@oslab-ASUSPRO-D840MB-M840MB:~/filesystem_hw1/hw1/ranWrite_direct$ ./ranWrite_direct.out  
ranWrite_direct exTime = 693525402 us
```

# Turn in

- Deadline  
2020/04/08 PM.11:59:59
- Upload to ilearning
- File name
  - ☐ HW1\_ID.zip (e.g. HW1\_7105056035.zip)
    - Source code
      - ☐ .c file
    - Word
- If you don't hand in your homework on time, your score will be deducted 10 points every day.



# Rules

- **No cheat work** is acceptable
  - You get zero if you copy other people's version.
- **Only single job** is accepted



# TA

- Name : Cheng-Chia, Chang
- Email : s9001055@gmail.com
  - Title format : FS HW#1 - [your name]
- Lab: OSNET(1001A)

# Reference

- Operating System Concepts, 10th Edition
- The Linux Programming Interface:  
A Linux and UNIX System Programming Handbook
- <https://systemprogrammingatntu.github.io/>